

## Optimization

Let  $f: \mathbb{R}^d \rightarrow \mathbb{R}$  be a scalar-valued function. ( $f(x_1, \dots, x_d$

Gradient:  $\nabla_x f(x) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_d} \end{bmatrix}$

Hessian:  $\nabla_x^2 f(x) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_d} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_d} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_d \partial x_1} & \dots & \dots & \frac{\partial^2 f}{\partial x_d^2} \end{bmatrix}$

⊛ In the case of

$f: \mathbb{R}^d \rightarrow \mathbb{R}^m$  being a vector-valued function, i.e.:

$$f(x) = (f_1(x_1, \dots, x_d), f_2(x_1, \dots, x_d), \dots, f_m(x_1, x_2, \dots, x_d))$$

Jacobian:  $\nabla_x f(x) = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \dots & \frac{\partial f_m}{\partial x_1} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_1}{\partial x_d} & \dots & \frac{\partial f_m}{\partial x_d} \end{bmatrix}$

---

Setup: Given a model with  $\overbrace{\vartheta \in \mathbb{R}^d}^{\text{parameters}}$ , i.e.  $\vartheta = (\vartheta_1, \vartheta_2, \dots, \vartheta_d)$

and a loss/likelihood  $L(\vartheta)$ , then our goal is to identify a set of parameters  $\vartheta^*$  such that:





$$\vartheta^* = \arg \min_{\vartheta} L(\vartheta)$$

We need to identify the critical points for which:

$$\nabla_{\theta} L(\theta) = 0$$

$$L: \mathbb{R}^d \rightarrow \mathbb{R}$$

This condition is true for: (i) minima, (ii) maxima, (iii) saddle point



### Gradient descent:


We want to minimize  $L(\theta)$ , i.e.  $\theta^* = \operatorname{argmin}_{\theta} L(\theta)$

Pick an initial guess  $\theta_0$  and use update rule:

Update Rule:  $\theta_{n+1} = \theta_n - \eta \nabla_{\theta} L(\theta_n)$

discretize version  $\rightarrow$  Forward Euler

$\frac{d\theta}{dt} = -\nabla_{\theta} L(\theta)$  gradient flow system



• This is a first-order method as it relies on a linear approximation of  $L(\theta)$  around some point  $\theta_n$ .

• It is guaranteed to converge to a critical point (e.g. local min/max) of  $L(\theta)$ , assuming that  $\eta$  is properly chosen.

$\eta$ : step-size/learning rate (user need to adjust)

### Newton's algorithm:

Let's use a Taylor expansion of  $L(\theta)$  around  $\theta_n$ :

$$\begin{aligned}
 \underset{1 \times 1}{L(\theta)} &\approx \underset{1 \times d}{L(\theta_n)} + \underset{d \times 1}{g_n^T} (\underset{d \times d}{\theta - \theta_n}) + \frac{1}{2} \underset{1 \times d}{(\theta - \theta_n)^T} \underset{d \times d}{H_n} (\underset{d \times 1}{\theta - \theta_n}) \quad \left. \begin{array}{l} \text{where:} \\ g_n := \nabla_{\theta} L(\theta_n) \\ H_n := \nabla_{\theta}^2 L(\theta_n) \end{array} \right\} \\
 &\approx L(\theta_n) + g_n^T (\theta - \theta_n) + \frac{1}{2} [\theta^T H_n \theta - 2\theta^T H_n \theta_n + \theta_n^T H_n \theta_n]
 \end{aligned}$$

Let us now find critical points:


$$\nabla_{\theta} L(\theta) = 0 \Rightarrow 0 + g_n^T + H_n \theta - H_n \theta_n = 0$$

$$\Rightarrow \theta = \theta_n - H_n^{-1} g_n^T$$

critical point of the Taylor approximation to  $L(\theta)$   
quadratic

Newton

update rule:  $\theta_{n+1} = \theta_n - H_n^{-1} g_n^T$



- No need for tuning parameters, but extra cost for computing and inverting the Hessian.

- 2<sup>nd</sup>-order algorithm, faster convergence vs gradient descent

- utilizes the underlying geometry by exploiting curvature information

Example: Linear regression

Recall the loss function for MLE estimation:

$$L(\theta) := -\log p(y | \underline{w}^T x, \sigma^2), \quad \theta := \{w, \sigma^2\}_{d \times 1, 1 \times 1}$$

Gradient:  $\nabla_w L(w) = -X^T y + X^T X w$

Hessian:  $\nabla_w^2 L(w) = X^T X$

Gradient descent:  $w_{n+1} = w_n - \eta [-X^T y + X^T X w_n]$

Newton:  $w_{n+1} = w_n - (X^T X)^{-1} [-X^T X w_n + X^T y]$

\* Remark : Similar updates can be formulated for  $\sigma^2$  !

### Limitations :

- Gradient descent converges slowly. Choosing  $\eta$  is an art.

① Scalability to big data.

evaluating the loss  
over a big data-set  
is expensive.

$$\theta_{n+1} = \theta_n - \eta \nabla_{\theta} L(\theta_n)$$

typically:  $\propto \sum_{i=1}^n (y_i - w^T x_i)$

- Exact Hessians are often very hard to compute  $\xrightarrow{\text{approx}}$  Quasi-Newton methods.

---

## Stochastic gradient descent (SGD)

In many ML applications the loss functions factorize across data points, i.e. they can be written as a summation over individual data points:  $L(\theta) \propto \sum_{i=1}^n L_i(\theta)$  (e.g. linear regression)

(this follows from assuming an i.i.d likelihood).

In this case, a standard "full batch" gradient descent approach would take the form:  $\theta_{n+1} = \theta_n - \eta \underbrace{\nabla_{\theta} L(\theta_n)}_{\text{true gradient}} = \theta_n - \eta \underbrace{\sum_{i=1}^n \nabla_{\theta} L_i(\theta_n)}_{\text{true gradient}}$

In stochastic gradient descent,

the true gradient is approximated using a single example:

$$\theta_{n+1} = \theta_n - \eta \underbrace{\nabla_{\theta} L_i(\theta_n)}_{\substack{\text{very noisy} \\ \text{approx. to} \\ \text{the true gradient}}}, \quad \begin{array}{l} i \text{ is chosen at random} \\ \text{at each iteration of} \\ \text{SGD algorithm} \end{array}$$

A compromise between these two extremes is to approximate the true gradient over a "mini-batch" of data:

$$\theta_{n+1} = \theta_n - \eta \frac{1}{M} \sum_{i=1}^M \nabla_{\theta} L_i(\theta_n) \quad M \ll n$$

$$u_{n+1} = u_n + \sum_{i=1}^m \dots$$

At every iteration we randomly pick a mini-batch of data by sub-sampling our full data-set.

- A complete looping cycle over the entire data-set is called an "epoch".

## Modern variants of SGD

- SGD with momentum:

$$u_n: \text{momentum variable} \quad \begin{cases} u_{n+1} = \gamma u_n + \eta \nabla_g \mathcal{L}(\theta_n) , & u_0 = 0 \\ \theta_{n+1} = \theta_n - u_{n+1} \end{cases} \quad \begin{cases} \eta: \text{learning rate} \\ \gamma \begin{cases} 0 = \text{SGD} \\ 0.9 \text{ is a typical value in prac.} \end{cases} \end{cases}$$

- Nesterov accelerated gradient method (NAG):

$$\begin{cases} u_{n+1} = \gamma u_n + \eta \nabla_g \mathcal{L}(\theta_n - \gamma u_n) \\ \theta_{n+1} = \theta_n - u_{n+1} \end{cases}$$

## Adaptive learning rate approaches:

### 1.) RMS prop:

$\mathbb{E}[g^2]_n := \text{average of the square gradients.}$

$\nwarrow \mathbb{E}[(\nabla_g \mathcal{L}(\theta_n))^2]_{dx \perp}$

running avg of the sq. gradients  $\nwarrow$

$$\begin{cases} \mathbb{E}[g^2]_{n+1} = \gamma \mathbb{E}[g^2]_n + (1-\gamma) g_n^2 \\ \theta_{n+1} = \theta_n - \frac{\eta}{\sqrt{\mathbb{E}[g^2]_{n+1} + \epsilon}} g_n \end{cases} \quad \begin{matrix} \text{typically } \gamma = 0.9 \\ \eta = 10^{-3} \end{matrix}$$

$\epsilon = 10^{-8}$

### 2.) Adam (adaptive moment estimation):

$$\begin{cases} \rightarrow m_{n+1} = \beta_1 m_n + (1-\beta_1) g_n \rightarrow \text{running average estimator of the expected gradient} \\ \rightarrow v_{n+1} = \beta_2 v_n + (1-\beta_2) g_n^2 \rightarrow \text{running average estimator of the variance of the } g \end{cases}$$

$m, v$  are usually initialized to be zero

∴ we typically initialize  $\theta$  to be zero, hence the estimates are biased towards zero. To counteract this bias, we can consider a correction:

$$\hat{m}_n = \frac{m_n}{1 - \beta_1^n}, \quad \hat{v}_n = \frac{v_n}{1 - \beta_2^n}$$

$$\text{Adam update rule} \quad \theta_{n+1} = \theta_n - \underbrace{\frac{\eta}{\sqrt{\hat{v}_{n+1} + \epsilon}}}_{d \times 1} \hat{m}_{n+1}$$