


Convolutional Neural Networks

- They are also made up of "neuron", and they have differentiable weights and biases.
- Each neuron receives some inputs, performs a dot product, (optionally) followed by a non-linearity.
- Fwd pass and loss are fully differentiable.

CNNs: Are tailored to address the unfavorable complexity of MLP, for high-dimensional inputs/outputs, by making the explicit assumption that our data "lives" on grid (e.g. images, time-series).


This assumption allows us to encode certain structure in our NN arch:



\Rightarrow Conv Nets (CNNs) are neural networks that use convolution instead of matrix multiplication/projection, in at least one of their layers

Convolution: Definition in 1D

Assume a time-series $x(t)$ that is sampled at regular intervals.



To filter out the noise we will compute some weighted average of the measurements. To do so, we will choose a weighting function $w(s)$, and obtain the smoothed signal as:

$$S(t) = \int_{-\infty}^{\infty} x(s) w(t-s) ds = \int_{-\infty}^{\infty} \dots$$

$$S(t) = \int_{-\infty}^{\infty} w(s) x(t-s) ds$$

or $S(t) = (x * w)(t)$

convolution is a commutative linear operation.

Discrete convolution in 1D:

$$S(t) = (x * w)(t) = \sum_{s=-\infty}^{+\infty} x(s) w(t-s) = \sum_{s=-\infty}^{+\infty} w(s) x(t-s)$$

→ x :	<u>input</u> $[x_1 x_2 x_3 x_4]$	}	$S_1 = w_1 x_1 + w_2 x_2$
→ w :	<u>conv. kernel</u> $[w_1 w_2]$		$S_2 = w_1 x_2 + w_2 x_3$
→ S :	<u>feature map</u> $[S_1 S_2 S_3]$		$S_3 = w_1 x_3 + w_2 x_4$

Discrete convolution for a 2D image ($m \times n$ pixels)

$$S(i, j) = \sum_m \sum_n I(m, n) K(i-m, j-n) = \sum_m \sum_n K(m, n) I(i-m, j-n)$$