

ENM 360: Introduction to Data-driven Modeling

Lecture #23: Sampling methods

Paris Perdikaris
November 17, 2020



The Metropolis algorithm

- ▶ Choose a symmetric proposal matrix Q . So, $Q_{ab} = Q_{ba}$.
- ▶ Initialize $x_0 \in X$.
- ▶ for $i \in 0, 1, 2, \dots, n - 1$:
 - ▶ Sample proposal x from $Q(x_i, x)$ if x is discrete, otherwise, $p(x | x_i)$.
 - ▶ Sample r from $\text{Uniform}(0, 1)$.
 - ▶ If

$$r < \frac{\tilde{\pi}(x)}{\tilde{\pi}(x_i)},$$

accept and $x_{i+1} = x$.

- ▶ Otherwise, reject and $x_{i+1} = x_i$.

Output: x_0, x_1, \dots, x_n

- ▶ Symmetric proposals include:


$$J(\theta^* | \theta^{(s)}) = \text{Uniform}(\theta^{(s)} - \delta, \theta^{(s)} + \delta)$$

and


$$J(\theta^* | \theta^{(s)}) = \text{Normal}(\theta^{(s)}, \delta^2).$$

The Metropolis algorithm

Markov chains



Posterior density



The Metropolis algorithm for Bayesian inference

Goal: We want to sample from

$$p(\theta \mid y) = \frac{f(y \mid \theta)\pi(\theta)}{m(y)}.$$

Typically, we don't know $m(y)$.

The notation is a bit more complicated, but the set up is the same.

We'll approach it a bit differently, but the idea is exactly the same.

We know $\pi(\theta)$ and $f(y \mid \theta)$, so we can draw samples from these.

Our notation here will be that we assume parameter values $\theta_1, \theta_2, \dots, \theta_s$ which are drawn from $\pi(\theta)$.

We assume a new parameter value comes in that is θ^* .

The Metropolis algorithm for Bayesian inference

The Metropolis algorithm proceeds as follows:

1. Sample $\theta^* \sim J(\theta \mid \theta^{(s)})$.
2. Compute the acceptance ratio (r):

$$r = \frac{p(\theta^* | y)}{p(\theta^{(s)} | y)} = \frac{p(y \mid \theta^*)p(\theta^*)}{p(y \mid \theta^{(s)})p(\theta^{(s)})}.$$

3. Let

$$\theta^{(s+1)} = \begin{cases} \theta^* & \text{with prob } \min(r, 1) \\ \theta^{(s)} & \text{otherwise.} \end{cases}$$

Remark: Step 3 can be accomplished by sampling $u \sim \text{Uniform}(0, 1)$ and setting $\theta^{(s+1)} = \theta^*$ if $u < r$ and setting $\theta^{(s+1)} = \theta^{(s)}$ otherwise.

Bayesian linear regression

$$y_i \sim \mathcal{N}(\beta_0 + \beta_1 x_i, 1/\tau)$$

or equivalently

$$y_i = \beta_0 + \beta_1 x_i + \epsilon, \quad \epsilon \sim \mathcal{N}(0, 1/\tau)$$

The likelihood for this model may be written as the product over N iid observations

$$L(y_1, \dots, y_N, x_1, \dots, x_N | \beta_0, \beta_1, \tau) = \prod_{i=1}^N \mathcal{N}(\beta_0 + \beta_1 x_i, 1/\tau)$$

Priors on the model parameters:

$$\beta_0 \sim \mathcal{N}(\mu_0, 1/\tau_0)$$

$$\beta_1 \sim \mathcal{N}(\mu_1, 1/\tau_1)$$

$$\tau \sim \text{Gamma}(\alpha, \beta)$$

Gibbs sampling

Gibbs sampling works as follows: suppose we have two parameters θ_1 and θ_2 and some data x .

Our goal is to find the posterior distribution of $p(\theta_1, \theta_2 || x)$.

Gibbs sampling algorithm:

1. Pick some initial $\theta_2^{(i)}$.
2. Sample $\theta_1^{(i+1)} \sim p(\theta_1 || \theta_2^{(i)}, x)$
3. Sample $\theta_2^{(i+1)} \sim p(\theta_2 || \theta_1^{(i+1)}, x)$

Then increment i and repeat K times to draw K samples.

This is equivalent to sampling new values for a given variable *while holding all others constant*.

The general approach to deriving an update for a variable is

1. Write down the posterior conditional density in log-form
2. Throw away all terms that don't depend on the current sampling variable
3. Pretend this is the density for your variable of interest and all other variables are fixed. What distribution does the log-density remind you of?
4. That's your conditional sampling density!

Pros: No parameters need to be tuned (e.g. vs MCMC that needs a proposal distribution)

Cons: It might be hard to analytically derive the conditional distributions.